AST: A library for modelling and manipulating coordinate systems[☆]

David S. Berry[a,b,*], Rodney F. Warren-Smith[c], Tim Jenness[a,d]

[a]*Joint Astronomy Centre, 660 N. Aʻohōkū Place, Hilo, HI 96720, USA*
[b]*East Asian Observatory, 660 N. Aʻohōkū Place, Hilo, HI 96720, USA*
[c]*RAL Space, STFC Rutherford Appleton Laboratory, Harwell Oxford, Didcot, Oxfordshire OX11 0QX, UK*
[d]*LSST Project Office, 933 N. Cherry Ave, Tucson, AZ 85721, USA*

Abstract

In view of increased interest in object-oriented systems for describing coordinate information, we present a description of the data model used by the Starlink AST library. AST provides a comprehensive range of facilities forattaching world co-ordinate systems to astronomical data, for retrievingand interpreting that information in a variety of formats, includingFITS-WCS. AST is a mature system that has been in use for more than 20 years, and may consequently be useful as a means of informing development of similar systems in the future.

Keywords:
WCS, data models, Starlink

1. Introduction

The Starlink AST library (Warren-Smith and Berry, 2013, ascl:1404.016) provides a generalised scheme for modelling, manipulating and storing inter-related coordinate systems. Whilst written in C, it has bindings for several other languages including Python, Java, Perl and Fortran. It has specialised support for many of the coordinate systems and projections commonly used to describe astronomical World Coordinate Systems (WCS), including all those described by the FITS-WCS standard, plus various popular distortion schemes currently in use. However, it is not limited to WCS, and may be used in any situation requiring transformation between different coordinate systems.

Unlike FITS-WCS, which supports only a relatively small set of prescribed transformation recipes reflecting the coordinate transformations within an optical telescope, AST allows arbitrarily complex transformations to be constructed by combining simple atomic transformations. This allows a much wider range of transformations to be described than is possible using FITS-WCS, and so can accommodate a wider range of data storage forms without the need to re-grid the data.

AST was released internally in 1997 (Warren-Smith and Berry, 1998, included in "Twenty Years of ADASS"). Since then it has been in continuous use within the Starlink Software Collection (Currie et al., 2014, ascl:1110.012) and is also used by various other major astronomical software tools such as DS9 (Joye and Mandel, 2003, ascl:0003.002). Interest in flexible schemes for representing inter-related coordinate systems has increased recently[1], and so it seems an appropriate time to review the lessons learned from AST.

This paper first presents an account of the historical issues that drove the initial development of AST, together with the reasoning behind some of the design decisions, and then presents an over-view of the more important aspects of the data model used by AST.

2. Historical Perspective

2.1. Initial Problems

The first public release of the AST library was in 1998 (Draper, 1998; Lawden, 1998) but some of the underlying concepts date from the late 1980s, when the Starlink Project was designing its NDF data format for gridded astronomical data (Jenness et al., 2015). There was clearly a need to relate positions within gridded data, using coordinates based on pixel indices, to real-world positions on the sky, wavelengths in a spectrum and so on. These non-pixel coordinate systems are now generally known as world coordinate systems.

Calibration of spectra, for example, was commonly performed by fitting a polynomial to express wavelength as a function of pixel position and then either storing the polynomial coefficients, or tabulating the polynomial value at each pixel centre. While not completely general, the latter option was an acceptable solution and was adopted as part of the Starlink data format. An array giving the central wavelength at each pixel was stored as the `AXIS` component in the NDF data structure and did good service in spectroscopic applications. It was also

---

[1]For instance, in the discussions about possible successors to the FITS format (Mink et al., 2015), and within the Astropy project (Astropy Collaboration, 2013)

possible, in a simple minded way, to attach an `AXIS` array to each dimension of an image or any gridded data set of higher dimensionality. This allowed each of its axes to be calibrated in terms of world coordinates.

This approach was adequate if the axes represented independent quantities (like wavelength and position for a long-slit spectrum), but did not suffice if the axes were inter-independent. Unfortunately, in the common case of celestial coordinates (such as Right Ascension and Declination), the axes are almost always inter-dependent. This is because the sky is essentially spherical and its coordinates are therefore naturally curvilinear when projected into two dimensions. This inter-dependence is a common feature of world coordinate systems in practice, so a solution was clearly needed that addressed it properly.

The Flexible Image Transport System (FITS; Greisen and Calabretta, 1995; Wells et al., 1981), at that time, addressed the issue in a better but still rudimentary way. In essence, it stored a physical pixel size (*e.g.*, in seconds of arc), allowed for a linear scaling of an image (typically to allow for the position angle rotation of the telescope) and then projected it on to the celestial sphere using one of a defined set of map projections. This representation was clearly based on a model of a physical telescope and how it imaged an observed region of the sky in its focal plane.

While successfully accommodating the curvilinear nature of sky coordinates, this FITS approach was still limited in many ways. In essence, it defined a small set of functional forms (based on map projections) through which pixel coordinates could be mapped on to celestial coordinates and back again. However, if the actual relationship between pixel coordinates and world coordinates didn't correspond to one of these functional forms, then it wasn't possible to use FITS to store the coordinate information[2].

For instance, if astronomical instrumentation were to use a novel map projection, if arbitrary instrumental distortions were present or if the data were re-gridded into a non-physical space, then the FITS approach would fail. It also had limited support for high-accuracy astrometry, where the departure of the sky from a perfect sphere, for a variety of reasons, has to be taken into account. In addition, there are many other non-celestial world coordinate systems that one might use (involving energy, velocity, time, frequency, *etc.*) that no contemporary system could represent adequately.

Unfortunately, this list of limitations only scratches the surface of the problem as it was perceived at the time. Other considerations, such as the time-dependent relationship between non-inertial celestial coordinate systems, the dependence of apparent positions on the position and velocity of the observer (and also on the wavelength of observation and atmospheric conditions) and periodic revisions to the fundamental definitions of celestial coordinate and time systems would all have to be accommodated, as would numerous other issues specific to particular domains (celestial coordinates, time systems, radial

velocities, wavelength/energy, *etc.*). This was several years before the FITS community commenced work on what was eventually to become the current FITS-WCS standard.

2.2. TRANSFORM

In the late 1980s, no immediate and general solution to these problems could be seen. Recognising the limitations in the FITS approach, however, the Starlink Project decided to take a hard line and to omit completely any component dedicated to world coordinate systems from its new NDF data format. Instead, this *astrometry extension* (from which the name AST is derived) was to be added at a later date when a suitable solution had been formulated.

This decision was undoubtedly strongly influenced by Patrick Wallace's presence in the Project and the major work he had done on the SLALIB library (Wallace, 1994, ascl:1403.025) to encapsulate best-practice in astrometric calculations (and also in other domains such as time systems). Discussion within the Project rapidly convinced us that if we adopted the FITS approach as it existed at the time, we would cut ourselves off from the proper rigorous treatment of astrometric data that is needed for the highest accuracy.

Consequently, a pilot project was conducted to explore alternative approaches. The most important limitation of the FITS approach was felt to be the use of a fixed set of functional forms (map projections) each of which was associated with a small fixed set of parameters. This simplified storing the information in 80-character FITS *header cards*, but clearly the set of functional forms that might ultimately be needed was much larger than had been recognised. Adding new ones might become a never-ending project and that, in turn, raised the prospect of continually upgrading all software that had to read and process FITS headers and handle coordinate systems.

The alternative approach that we explored was to write an expression parser that would accept sets of arithmetic expressions similar to those used in Fortran and C, along with the usual set of mathematical functions. Together with a method of passing named parameter values into these expressions, this greatly increased the set of functional forms that could be represented. The expressions themselves (encoded as character strings) and the associated parameter values could easily be stored in astronomical data sets. Typically, one set of expressions would relate pixel coordinates to world coordinates (*e.g.*, sky coordinates) and a second, optional, set would define the inverse transformation. The expression syntax was powerful enough to represent a wide range of map projections plus many other transformations into alternative world coordinate systems.

A processing engine was also provided that could use the stored expression data to transform actual coordinate values.

A library implementing this, called TRANSFORM, was released in 1989 (Lawden, 1989; Warren-Smith, 1989). It stored its data (the expressions and parameters) in Starlink's Hierarchical Data Format (HDS; Jenness, 2015; Lupton, 1989; Warren-Smith et al., 2008) and was thus able to integrate with the Starlink NDF data format to attach arbitrary world coordinates to gridded astronomical data sets.

---

[2]Unless the data was first re-gridded into a form supported by FITS.

2.3. TRANSFORM Lessons

Ultimately, TRANSFORM turned out not to be a full solution to the WCS problem and did not become part of the NDF data format[3]. It was, however, used for two initially unforeseen purposes which turned out to be very significant:

1. Associating coordinate systems with plotting surfaces in a "graphics database" (see *e.g.*, Eaton and McIlwrath, 2013). This allowed plotting applications to store a coordinate system for (say) a graph plotted in logarithmic coordinates so that those coordinates could later be recovered from the position of a cursor. This demonstrated that plotting was a major application area for this type of technology, especially when using curvilinear coordinates such as Right Ascension and Declination which are notoriously difficult to handle properly with standard plotting software.

2. Transformation and combination of bulk image data using general arithmetic expressions (as an alternative to combining images using a manual sequence of add/subtract/multiply/divide and similar applications). This showed that (a) the approach could easily be efficient enough to handle large data sets and (b) the data values in an image were just another coordinate that could be transformed into different representations (logarithmic, different units, *etc.*) in much the same way as its axes.

With these insights, it was clear that the ideas behind TRANSFORM had potential, but some serious deficiencies had also emerged:

- Arithmetic expressions, while fairly general, could not easily cope with coordinate transformations that required iterative solution, nor with discontinuous transformations, nor with look-up tables or a variety of other computational techniques. While arithmetic expressions provided a valuable increase in the flexibility of coordinate transformations, clearly other classes were still needed.

- It was a major problem for the average writer of astronomical software to formulate the required coordinate expressions correctly even when dealing with quite simple sky coordinate systems. The core of this issue is that celestial coordinate systems are rather complex and a good deal of specialist knowledge is needed to formulate even simple cases correctly. Clearly a better solution would be to encapsulate this knowledge in the WCS software and provide a simpler API that dealt only in high-level concepts.

- For high accuracy work, further complex calculations arise. These are related, for example, to atmospheric refraction and special & general relativistic effects (like the observer's motion and the sun's gravity). These require the use of a dedicated library of astrometric functions and cannot in practice be handled by simple expressions. They also require additional data about the observing context

(time, position, velocity, wavelength, *etc.*) and any practical solution must define how these are stored and processed.

- TRANSFORM had no ability to store additional information about data axes, such as labels and units.

- It became clear that coordinate transformations frequently needed to be combined, for example by applying one transformation after another, and that this process was often inefficient. The key to better efficiency lay in knowing more about each transformation, like whether it was linear or had a variety of other properties. With this information it was possible to merge (or cancel out) consecutive transformations for better efficiency. TRANSFORM had a rudimentary system for encoding this information, but it was not really up to the task.

- Tying WCS software to a particular (Starlink) data system was a mistake and limited the uses to which it could be put. It would clearly be better if the data could be encoded (serialised) in alternative ways to make it data-system agnostic. The same agnosticism should also apply to other likely dependencies, like graphics systems and error reporting. The ability to implement these services in alternative ways would be especially important when designing graphical user interfaces that processed WCS information.

2.4. Developments in FITS WCS

At about the same time, the wider FITS community also came to recognise some of the limitations of WCS handling within FITS, and in 1992 work commenced on a new standard for storing WCS information within FITS files. However, in view of the "once FITS, always FITS" principle, that work consisted mainly in formalising and extending existing practices. So for instance, new keywords were defined to store the extra meta-data needed for a complete description of a celestial coordinate system, and new projection types were added, but the basic model remained unchanged. The new standard still required that the transformation be split into three components applied in series; an affine transformation that converts pixel coordinates into *intermediate world coordinates*, a spherical projection that converts these into *native spherical coordinates*, and a spherical rotation that converts these into the final world coordinate system.

In view of the decision to stay with this rigid and restrictive model, and the expected length of time needed to agree a new standard[4], the Starlink project decided in early 1996 to develop its own WCS system, informed by the earlier experiments with TRANSFORM, rather than adopt the new FITS standard.

2.5. AST Principles

One of the first decisions was to separate the representation of a coordinate system (that we called a Frame) from the computational recipe that transforms between coordinate systems (a

---

[3]Although it it was the precursor of the `MathMap` class in AST.

[4]An expectation that was justified when the standard was finally published in 2002.

Mapping). From the TRANSFORM experience we knew we would need multiple classes of both these data types, all of which would need to support the same basic operations, but each having its own specialisation. The correspondence with sub-classing in object-oriented (OO) programming was irresistible and the decision to use an OO design immediately followed.

This raised the issue of an implementation language. We planned to use the SLALIB library for astrometric calculations[5]. This had been developed with extreme portability in mind and had recently been re-written in ANSI C. We didn't want to compromise this portability, so decided also to work in strict ANSI C and to minimise software dependencies as much as possible. This meant providing portable interfaces to facilities that were intrinsically less portable, such as data file access, plotting, error reporting, *etc.* and providing simple implementations that users could re-write if necessary.

Deciding to write an OO system in a non-OO language took considerable thought. We needed to provide a Fortran-callable interface but, at the time, the portability of C++ code was quite limited if one needed to call it from Fortran, so that route was unattractive[6]. Eventually, we were guided by the approach described by Holub (1992) for handling objects in C and were able to hide the detail from users using pre-processor macros.

One consequence of this is that users cannot easily create new sub-classes from AST objects without learning the internal conventions that it uses. At the time, this was seen as something of an advantage. The library is intended for data interchange and creating new sub-classes would inevitably allow persistent objects to be created that other users could not access. However, with hind-sight a more open architecture may have encouraged involvement from a wider user-base. [7].

As noted previously, we wanted the AST API to deal in high-level concepts and to hide as much specialist detail as possible from the user. This principle arose from considering the complex calculations involved in handling celestial coordinates and time systems. However, we soon realised that two other areas were similarly complex and could benefit from the same approach.

The first area was graphics. Plotting in curvilinear coordinates is a complicated business if one wants to handle all the corner cases correctly. Plotting and labelling celestial coordinate axes, for example, presents many problems; especially near the poles of an all-sky projection. It is made even harder if the projection contains discontinuities. But the high-level concepts involved in such plotting (coordinate systems and the mappings between them) are such a natural fit with other AST concepts that it seemed obvious to implement a class of coordinate system that is specialised for graphics. The high-level operations it supports would then hide the details of the complex and generalised plotting algorithms involved.

The second area is an aspect of data storage – namely the handling of FITS header cards. While the AST library could provide ways to serialise its own data transparently, possibly in multiple ways, it also needed to inter-operate with FITS. WCS data in FITS data files is stored in a series of 80-character header cards and, over the years, the number of different ways the information can be stored in these headers has grown. The complexity involved is now considerable (see *e.g.*, Thomas et al., 2015). Again, detailed specialist knowledge is needed to extract this information reliably and to write it back (possibly modified) in a form that gives other FITS-handling software a chance of using it while not conflicting with the many other FITS headers typically present.

For a user of the AST library, we wanted this process of accessing FITS headers to appear as much like a simple read/write operation as possible, with all the implementation details hidden. This requirement arose from more than simply ease-of-use. FITS header conventions (many of them informal) are in constant flux and if these details are embedded in applications programs, those programs must constantly receive attention if they are to remain up to date. Embedding all these details in the AST library allows the problem to be addressed in one place and by someone with the necessary expertise.

FITS header handling has proved one of the most complex areas to tame in AST. But introducing the concept of a *destructive read* (which reads WCS data from FITS headers and simultaneously deletes the relevant headers) has made it possible to write applications with very little code that have completely general handling of FITS WCS headers (see section 4.1.1).

3. The AST Data Model

The most basic principle behind the AST data model is a clear distinction between a *transformation* and a *coordinate system*.

A *transformation* is a mathematical recipe for converting a numerical input vector into a corresponding numerical output vector. The transformation itself has no knowledge of what the values within these vectors represent, other than that each one constitutes a position within some unspecified N-dimensional space[8].

A *coordinate system* is a collection of meta-data describing a set of one or more axes. This will include:

- the number of axes (i.e., dimensionality of the space)

- the physical quantity described by each axis

- the units used by each axis

- the geometry of the space that they describe (flat, spherical, *etc*)

- the nature of the coordinate system (Cartesian, polar, *etc*)

---

[5]A later version of AST eventually replaced SLALIB with SOFA and PAL.

[6]Further, there was no official C++ standard available at the time - the first international standard for C++ (ISO/IEC 14882:1998) was published in 1998.

[7]For Mappings, where this issue is most relevant, the problem has been mitigated in a controlled way by the IntraMap class that allows separately-compiled code to be imported into the library.

[8]The dimensionality of the output space need not be the same as that of the input space.

- other meta-data that may be needed to specify the coordinate system fully.

These two concepts are encapsulated within two separate classes in the AST data model: the *Mapping* class and the *Frame* class. This separation underlines the fundamental difference between the two main requirements of any coordinate handling system:

1. knowing *how* to convert numerical positions from one coordinate system to another.
2. knowing *what* those coordinate systems represent.

As an example of the practical consequences of this distinction, the *pixel size* of a typical 2-dimensional image of the sky is *not* considered to be a property of the (RA,Dec) Frame, since it is determined by the nature of the transformation from pixel coordinates to (RA,Dec) coordinates, rather than being an intrinsic property of the (RA,Dec) coordinate system itself.

The two base classes, *Mapping* and *Frame*, are extended to create a wide variety of sub-classes, each of which describes a specific form of transformation or coordinate system. New sub-classes can be added as required and slot naturally into the existing infra-structure provided by the rest of AST.

In addition, this separation into two "orthogonal" classes makes it easy to create complex compound objects from simple component objects. For instance, multiple Mappings can be combined into a new object, and the resulting object will itself be a Mapping. Likewise, multiple Frames can be combined into a new object, and the resulting object will itself be a Frame[9].

However, at some point these two classes need to be brought together to provide a complete description of a set of related coordinate systems. The *FrameSet* class is used for this purpose. A *FrameSet* encapsulates a collection of two or more Frames, with the Mappings that describe the transformation between the corresponding coordinate systems. The simplest FrameSet contains two Frames, together with a single Mapping that describes how to convert positions between these two Frames (see Fig. 1).


Figure 1: A FrameSet that describes two coordinate systems and the transformations between them. The Mapping's *forward* transformation transforms positions in *Frame 1* to the corresponding position in *Frame 2*. The Mapping's *inverse* transformation transforms positions in *Frame 2* to the corresponding position in *Frame 1*.

More complex FrameSets can be created that describe the relationships between multiple Frames in the form of a tree structure (see Fig. 2).

3.1. Transformations and Mappings

Within AST, most *Mappings* encapsulate two transformations - one is designated as the *forward* transformation and the


Figure 2: A FrameSet that describes five inter-related coordinate systems and the transformations between them.

other as the *inverse* transformation. When a Mapping is used to transform a set of positions, the caller must indicate if the forward or inverse transformation is to be used. The *forward transformation* converts positions within the input space of the Mapping into corresponding positions within the output space, and the *inverse transformation* converts positions within the output space of the Mapping into corresponding positions within the input space. A Mapping can be *inverted*, which results in the two transformations being swapped.

For most classes of Mapping, the inverse transformation is a genuine mathematical inverse of the forward transformation. However, this is not an absolute requirement, and there are a few classes of Mapping where this is not the case (for instance the *PermMap* class, when the axes of the output space are a permuted subset of the axes of the input space). In addition, the Mapping class does not require that *both* transformations are defined. For instance, the *MatrixMap* class, which multiplies each input vector by a specified matrix to create the output vector, will only have an inverse transformation if the matrix is square and invertible.

3.1.1. Atomic Mappings Provided by AST

AST provides many classes of Mapping that implement a wide range of different transformations. Most of these are *atomic* Mappings that implement a specific numerical transformation and, if possible, its inverse. But some are *compound* Mappings that combine together other Mappings (atomic or compound) in various ways to create a more complex Mapping. A compound Mappings does not define its own transformations, but instead inherits the transformations of the individual component Mappings which it encapsulates.

The most significant atomic Mapping classes are:

UnitMap: Copy positions from input to output without any change.

WinMap: Transform positions by scaling and shifting each axis.

ZoomMap: Transform positions by zooming all axes about the origin.

---

[9]For instance, a 2-dimensional (RA,Dec) Frame can be combined with a 1-dimensional wavelength Frame to create a 3-dimensional (RA,Dec,Wavelength) Frame.

ShiftMap: Translate positions by adding an offset to each axis.

MatrixMap: Transform position vectors by multiplying by a matrix.

PolyMap: A general N-dimensional polynomial transformation.

PermMap: Transform positions by permuting and selecting axes.

LutMap: Transform 1-dimensional coordinates using a look-up table.

MathMap: Transform coordinates using general algebraic mathematical expressions.

WcsMap: Implements a wide range of spherical projections.

SlaMap: Transform positions between various celestial coordinate systems.

SpecMap: Transforms positions between various spectral coordinate systems.

SphMap: Map 3-d Cartesian to 2-d spherical coordinates.

TimeMap: Transform positions between various time coordinate systems.

PcdMap: Apply 2-dimensional pincushion/barrel distortion.

DssMap: Transform positions using Digitised Sky Survey plate solutions.

GrismMap: Models the spectral dispersion produced by a grism.

IntraMap: Transform positions using an externally supplied transformation function.

SelectorMap: Locates positions within a set of Regions (see section 3.1.2).

All classes of Mapping are immutable. That is, a Mapping cannot be changed once it has been created. This is unlike Frame and FrameSet objects, which *can* be changed.

3.1.2. Compound Mappings

The two most significant compound Mapping classes are:

CmpMap: The CmpMap class is a subclass of Mapping that encapsulates two other Mappings either in series or in parallel. Either or both of the two encapsulated Mappings can itself be a CmpMap, allowing arbitrarily complex Mappings to be created.

In a *series* CmpMap, each input position is transformed by the first component Mapping, and the output from that Mapping is then transformed by the second component Mapping. Consequently, the output dimensionality of the first Mapping must be the same as the input dimensionality of the second Mapping.

In a *parallel* CmpMap, the input space is split into two sub-spaces. The first component Mapping is used to transform the axis values corresponding to the first subspace, and the second component Mapping is used to transform the axis values corresponding to the second subspace. Thus the input dimensionality of the CmpMap is equal to the sum of the input dimensionalities of the two component Mappings, and the output dimensionality of the CmpMap is equal to the sum of the output dimensionalities of the two component Mappings.

SwitchMap: The SwitchMap class is a subclass of Mapping that allows a different transformation to be used for different regions within the input space.

Each SwitchMap encapsulates any number of other Mappings, known as "route" Mappings, and one "selector" Mapping. All of these Mappings must have the same input dimensionality, and all the route Mappings must have the same output dimensionality. The selector Mapping must have a one-dimensional output space.

Each input position supplied to the SwitchMap is first transformed by the selector Mapping. The scalar output from the selector Mapping is used to index into the list of route mappings. The selected route mapping is then used to transform the input position to generate the output position returned by the SwitchMap.

There is a specialised subclass of Mapping, the SelectorMap class, that is designed specifically to fulfil the role of the selector Mapping within a SwitchMap, but in principle *any* suitable form of Mapping may be used. The SelectorMap class encapsulates several Regions (see section 3.4) and returns an output value that indicates which of the Regions (if any) contained the input value. Thus, the SwitchMap would typically contain one route Mapping for each of the Regions contained within the SelectorMap.

3.1.3. Simplification

There are a wide range of possible transformations that could potentially be applied to a data set during analysis. These including simple things such as rotation, scaling, shear, *etc.*, but could in principle include more complex transformations such as re-projection, dis-continuous "patchwork" transformations, or even transformation using a general algebraic expression. A coordinate handling system should make it possible for a user to apply an arbitrary set of such transformations in series to a data set, without losing track of the coordinates of each data point. With a prescriptive scheme such as FITS-WCS this would require each transformation to locate the appropriate component of the FITS-WCS pixel to world coordinate mapping, and modify the corresponding headers in a suitable way. This is often a difficult, if not impossible, task. Within AST, the chaining of transformations is accomplished simply by creating a Mapping that describes each new transformation and concatenating it with the existing pixel to world coordinate mapping by creating a new CmpMap (see section 3.1.2).

However, by itself this can lead to Mappings that becoming increasingly complex as transformations are stacked on top of each other. This is a problem because it leads to:

1. slower evaluation of the total transformation,
2. less accurate evaluation of the total transformation, and
3. more room being needed for storage.

To avoid this, the Mapping class provides a method that takes a potentially complex Mapping and simplifies it as far as possible. Doing such simplification in a general and effective manner is one of the most difficult challenges faced by the AST model, but experience has shown that the current scheme handles most cases sufficiently well. The steps involved in simplification depend on the nature of the component Mappings in the total CmpMap. Each class of Mapping provides its own rules that indicate when and how it can be simplified, or combined with an adjacent Mapping in the chain. To illustrate the principle, some of the simplest examples include:

1. any Mapping can be combined with its own inverse to create a UnitMap,
2. UnitMaps can be removed entirely,
3. adjacent MatrixMaps in series can be combined using matrix multiplication to create a single MatrixMap,
4. adjacent MatrixMaps in parallel can be combined to create a single MatrixMap of higher dimensionality (filling the off-diagonal quadrants with zeros),
5. adjacent ShiftMaps can be combined to form a single ShiftMap (either in series or in parallel).

The whole simplification process is managed by the CmpMap class - it expands the compound Mapping into a list of atomic Mappings to be applied in series or parallel, and then for each Mapping in the list, invokes that Mapping's protected `astMapMerge` method. This method is supplied with the entire list of atomic Mappings, and determines if the nominated Mapping can be merged with any of its neighbours. If so, a new list of Mappings is returned containing the merged Mapping in place of the original mappings. Once all atomic Mappings in the CmpMap have been checked in this way, the same process is repeated again from the beginning in case any of the changes that have been made to the list allow further simplifications to be performed. This process is repeated until no further simplifications occur.

There are in general multiple ways in which a list of Mappings (either series or parallel) can be simplified. Since each class of Mapping has its own priorities about how to merge itself with its neighbours, it is possible for the simplification process to enter an infinite loop in which neighbouring Mappings disagree about the best way to simplify the Mapping list. When a particular Mapping is asked to merge with its neighbours, it may make changes to the Mapping list that are then undone when a neighbouring Mapping is asked to merge with its neighbours. The simplification method takes care to spot such loops and to assign priority to one or the other of the conflicting Mappings.

3.1.4. Missing or Bad Axis Values

AST flags unknown or missing axis values using a special numerical value, known as the "bad" value. If an input position supplied to a Mapping contains one or more bad axis values, then in general all output axis values will be bad[10].

Mappings may also generate bad output axis values if the input position corresponds to a singularity in the transformation, or is outside the region in which the transformation is defined.

3.2. Frames and Domains

As described in section 3, each instance of the Frame class contains all the meta-data necessary to give a complete description of a particular coordinate system. Each Frame is associated with a specific *domain* as explained in the next section.

Unlike Mappings, the properties of a Frame can be changed at any time.

3.2.1. What is a Domain?

AST uses the word *domain* to refer to a physical (or abstract) space such as "time", "the sky", "the electro-magnetic spectrum", "the focal plane", "a pixel array". Points within such a space can in general be described using any one of several coordinate systems. For instance, any position on the sky can be described using ICRS coordinates, Galactic coordinates, *etc*. Similarly, positions in the eletro-magnetic spectrum can be described using frequency, wavelength, velocity, *etc*.

Each subclass of Frame represents a specific domain with a domain name (a string) which is the same for all instances of the class. However, each instance of the class represents just one particular coordinate system within that domain. In general, each Frame sub-class will also encapsulate all the information needed to create a Mapping between any pair of supported coordinate systems within its domain.

For example, the *SkyFrame* class represents the astronomical sky (its domain) and has the domain name "SKY". It knows about a range of celestial coordinate systems that can be used to describe positions on the sky. However, an instance of the SkyFrame class represents only one of these celestial coordinate systems at a time. The particular one it represents is stored in its "system" attribute (an instance variable) which takes values such as "ICRS", "FK5", "Galactic", *etc*.

The SkyFrame class extends the Frame class by incorporating various other items of metadata necessary to perform conversion between the supported celestial coordinate systems, the main items being the epoch of observation and the reference equinox [11]. These are also instance variables and therefore specific to each SkyFrame instance. The coordinate system represented by an instance of the SkyFrame class can be changed at any time simply by assigning new values to any of the relevant attributes, as in the following Python examples:

```
import starlink.Ast as Ast
```

---

[10]One exception is that a parallel CmpMap may be able to generate non-bad values for some of its output axes.

[11]The base Frame class includes the observer's geodetic position.

```
# Create a Frame describing ICRS Right
# Ascension and Declination. This is the
# default system for the SkyFrame class.
frame1 = Ast.SkyFrame()

# Create a deep copy, and change the System
# attribute of the copy so that it represents
# Galactic longitude and latitude.
frame2 = frame1.copy()
frame2.System = "Galactic"
```

3.2.2. Coordinate Conversion within Domains

Given the two SkyFrames from the example in the previous section, it is possible to create a Mapping between them (*i.e.* from FK5 (RA,Dec) to Galactic (l,b)) based on the meta-data stored within the two SkyFrames. This Mapping can then be used to convert (RA,Dec) positions into equivalent (l,b) positions, or vice-versa. The process of creating this Mapping is implemented within the `astConvert` method of the Frame class. For instance:

```
my_frameset = frame1.convert( frame2 )
```

will, if possible, generate a Mapping from `frame1` to `frame2`. The returned Mapping is encapsulated within a FrameSet that also includes copies of `frame1` and `frame2`.

Likewise, the SpecFrame class encapsulates all the information needed to create Mappings between any pair of supported spectral coordinate system, accounting for rest frequency, standard of rest, celestial reference position, *etc.*, as in the following example:

```
import starlink.Ast as Ast

#   Create a Frame describing helio-centric
#   Radio Velocity in units of km/s, with a rest
#   frequency of 345.8 GHz.
frame1 = Ast.SpecFrame( "System=VRAD,RestFreq=345.8" )

#   Create a deep copy, and change the attributes
#   so that the copy describes frequency units of
#   "Hz" with respect to the kinematic Local Standard
#   of Rest.
frame2 = frame1.copy()
frame2.System = "FREQ"
frame2.Unit_1 = "Hz"
frame2.StdOfRest = "LSR"

#   Create a FrameSet that contains the Mapping
#   between these two spectral coordinate systems.
my_frameset = frame1.convert( frame2 )
```

The principle that each class of Frame contains all the meta-data and intelligence required to create a Mapping between any two coordinate systems within the Frame's domain, extends to compound Frames as well as atomic Frames, as described in the next section.

The base Frame class itself is slightly unusual in that it can be used to describe any generic domain, and is restricted to a single Cartesian coordinate system within that domain. The domain associated with a basic Frame is specified by the caller and can be any arbitrary string. Clearly, this restricts the usefulness of a basic Frame (compared to more specialised classes of Frame) in that it is not possible to include any knowledge about multiple coordinate systems given the arbitrary nature of the domain. The only exception is that the basic Frame class knows how to convert between different dimensionally equivalent units. Thus the implementation of the `astConvert` method provided by the

basic Frame class can generate a Mapping between two basic Frames if they have the same domain name, the same number of axes, and the axes have dimensionally equivalent units.

3.2.3. Atomic Frame Classes Provided by AST

AST provides several classes of Frame that describe different specialised domains. As with Mappings, these can be divided into *atomic* Frames that describe a single specific domain, and *compound* Frames that combine together other Frames (atomic or compound) to create a Frame describing a domain of higher dimensionality. The atomic Frame classes are:

Frame: An arbitrary N-dimensional domain with a single Cartesian coordinate system.

FluxFrame: A 1-dimensional domain describing several forms of flux measurement systems (all measured at a single spectral position). In this case, the "axis value" represents a flux value.

SkyFrame: A 2-dimensional domain describing several celestial coordinate systems.

SpecFrame: A 1-dimensional domain describing several spectral coordinate systems.

DSBSpecFrame: Extends the *SpecFrame* class to describe dual sideband spectral coordinate systems.

TimeFrame: A 1-dimensional domain describing several time coordinate systems.

3.2.4. Compound Frames

The following compound Frame classes are provided:

CmpFrame: The *CmpFrame* (compound Frame) class describes a coordinate system that combines the axes from two other Frames, in any order. The name of the domain associated with a CmpFrame is constructed automatically from the domain names of the two component Frames. For instance if a CmpFrame contains a SkyFrame and a SpecFrame, then its domain name will be "SKY-SPECTRUM".

SpecFluxFrame : The *SpecFluxFrame* class combines a FluxFrame with a SpecFrame.

3.2.5. Coordinate Conversion within Compound Domains

The principle that Frame classes contain all the information necessary to convert between any coordinate systems within their domain also applies to compound Frames and the compound domains that they represent. This facility is accessed through the astConvert method, as with atomic Frames (see section 3.2.2).

With atomic Frames, conversion is only possible between instances of the same Frame class (i.e. Frames that represent alternative coordinate systems within the same physical domain). However, with a compound Frame containing (say) two subcomponents, it may be possible to find a conversion to either

or both of the two sub-Frames, so that a new range of possibilities opens up. For instance, a CmpFrame describing the 3-dimensional "SKY-SPECTRUM" domain might be matched with any of the following:

- Another CmpFrame describing the SKY-SPECTRUM domain.

- Another CmpFrame describing the SPECTRUM-SKY domain.

- A SkyFrame (i.e., a Frame describing the SKY domain).

- A SpecFrame (i.e., a Frame describing the SPECTRUM domain).

where a "match" means that conversion is possible. If the destination Frame has fewer axes than the source Frame, then the resulting Mapping will contain a PermMap - an atomic Mapping that permutes and selects a subset of its input axes.

Because CmpFrames can be nested arbitrarily deeply and their axes can be permuted, deciding whether conversion is possible can be non-trivial, but the facility has great power and can form the basis of a search function for coordinate systems. For instance, if a general purpose program reads the WCS from a data file of arbitrary dimensionality and wants to ask the question "can I determine the ICRS (RA,Dec) of each pixel position?" it can create a SkyFrame describing ICRS (RA,dec) and then use the `astConvert` method to see if a Mapping can be created from the WCS Frame read from the data file to the *template* SkyFrame describing the required coordinate system. This will allow ICRS (RA,Dec) to be determined for any 2-dimensional data file calibrated in any of the supported celestial coordinate systems, and also for any multi-dimensional data file that contains a pair of celestial coordinate axes.

3.2.6. Other Important Frame Methods

A primary goal of AST is to support high-level software that can process arbitrary coordinate systems without needing to know their details. This simplifies the high-level algorithms themselves and also allows them to benefit from later additions to the AST library (*e.g.* new coordinate systems and domains) that did not exist when they were originally written. The object-oriented design of AST is key to this feature and it particularly affects the Frame classes.

To this end, the base Frame class declares a range of methods that provide a generic interface for handling coordinate systems. These methods are then typically over-ridden by sub-classes so that they exhibit the required more specialised behaviour when manipulated by generic high-level algorithms. Rather than detail the interface in full, we will illustrate the principle with a few examples here.

Encoding and decoding of axis values to/from text is a function that differs in detail between Frame sub-classes because of the differing conventions for formatting, say, simple floating point values, angles and times (the last two having a wide range of possible representations). The base Frame class defines a decoding interface that operates on arbitrary input text and converts it into axis values, regardless of the nature of the formatting involved or, indeed, which characters and delimiters are present. With suitable over-rides to the implementation in Frame sub-classes, this interface allows high-level software to read free-format input containing axis data for arbitrary Frames, including compound Frames, using a simple algorithm.

A matching encoding interface formats axis values under the control of a "Format" attribute string whose syntax can be over-ridden in Frame sub-classes to suite the formatting required. For example, the C-like "%.4f" might describe the formatting for a simple floating point axis, while the more specialised "hh:mm:ss" could be used for a Right Ascension axis. Being Frame-specific, however, these strings need to be set by specialised software, although defaults are provided. For a more generic interface, a "Digits" attribute is available that controls the number of significant digits shown and applies to any class of Frame. This finds particular use during graphical operations (see section 5.1).

A related method defined by the base Frame class will wrap axis values into a standard range (*e.g.* 0 to 360 degrees) where appropriate. This can be useful when reading and writing axis values, but also provides a simple way to characterise the different topologies of (say) flat, cylindrical and spherical spaces. A further method will return the distance between pairs of points, thereby defining the metric of the space that a Frame represents. The inverse operation is to return the coordinates of a point which is offset a given distance from a starting position in a defined direction.

These methods, suitably over-ridden in each Frame sub-class, form the basis for many generic graphics operations such as drawing geodesic curves (which replace straight lines as the basic drawing element in AST) and coordinate axes. For example, to mark formatted numerical values on the possibly curved axis of a graph, the encoding function described above is used. The formatting precision ("Digits" attribute) is progressively increased until all the formatted label values become distinct. This ensures adequate, but not excessive precision regardless of the range of axis values present and the algorithm works regardless of the type of Frame involved.

3.3. FrameSets

Converting between different coordinate systems within the same domain is handled by Frame and its sub-classes. However, converting between different domains (*e.g.* between pixel coordinates and sky coordinates) is not possible in this way because there is no intrinsic relationship between the two domains unless extra information is supplied. This extra information typically describes experimental configurations; for example where a telescope was pointing (or the settings of a spectrograph). Given this additional detail, it becomes possible to tie the two domains together, so that conversion can be performed between a pixel-based coordinate system and any celestial coordinate system (or wavelength system in the case of a spectrum).

Linking domains together in this way essentially creates a "super domain" where otherwise unconnected coordinate systems are bound together in a similar relationship to coordinate

systems within a single domain. This facility is typically required when calibrating the coordinate systems of experimental data and is provided by the *FrameSet* class. It also finds use in graphics applications for attaching a variety of (typically curvilinear) coordinate systems to a plotting surface.

Each domain within this "super domain", is represented by a separate Frame. If a FrameSet is used, for instance, to provide a complete description of the WCS associated with a data array, then one of the Frames within the FrameSet will represent pixel coordinates, and the other Frames will represent a collection of alternative world coordinate systems. For instance, a FrameSet describing an image taken by a telescope may have three Frames describing pixel coordinates, focal plane coordinates and sky coordinates.

3.3.1. FrameSets as Tree Structures

A FrameSet represents a network of inter-related coordinate systems in the form of a tree-structure in which each node is a Frame, with the nodes being connected together by Mappings (see Fig. 2).

The FrameSet class provides a method that returns the Mapping between any two nominated Frames. This may involve concatenating several Mappings if the two Frames are not directly connected to each other. For instance, if asked to return the Mapping between Frame 1 and Frame 4 in Fig. 2, the method will retrieve Mappings 1 and 3 from the FrameSet, and combine them in series into a single *CmpMap* (compound Mapping).

The FrameSet class also provides a method that returns any nominated Frame from a FrameSet.

New Frames can be added to an existing FrameSet at any point in the tree structure. To do so, the caller must provide a Mapping that maps positions from an existing Frame in the FrameSet to the new Frame.

Frames can also be removed from a FrameSet. If the Frame is a leaf node in the tree structure, then it is simply removed, together with the Mapping that connects it to its parent Frame. If the Frame is *not* a leaf node, the Frame is removed but the Mapping that connects it to its parent Frame is retained so that its child Frames can still be reached.

3.3.2. Base and Current Frames

Two of the most common operations on FrameSets are 1) converting positions from one Frame to another, and 2) enquiring or using the properties of a Frame. Performing these two operations repeatedly would become tedious if each such operation involved separate calls to extract the required Mapping or Frame from the FrameSet. To avoid this, AST is implemented in such a way that the FrameSet class effectively inherits from both the Mapping class and the Frame class. This means that any Mapping method, or any Frame method, can also be used on a FrameSet.

The FrameSet class allows two nominated Frames to be flagged within each FrameSet. One is referred to as the "current Frame", and the other as the "base Frame"[12]. When used as a Frame, a FrameSet is equivalent to its current Frame. When used as a Mapping, a FrameSet is equivalent to the Mapping from its base Frame to its current Frame.

When attaching WCS information to other data a FrameSet is typically employed and, by convention, its base Frame is used to represent the intrinsic coordinate system associated with the data. For example, the base Frame may be identified with the pixel coordinate system of an image or with the native coordinates used to address a plotting surface. The other Frames in the FrameSet represent alternative coordinate systems with which a user may choose to work and the current Frame represents the currently-selected coordinate system.

For example, an application that locates objects in images may generate results in pixel coordinates, but can then transform them between the base and current Frame coordinate systems in order to present them to the user in the required form (possibly as celestial coordinates). Similarly, a cursor position read from a graphics device can be shown to the user not in device coordinates, but in a form they can more readily understand.

When AST reads WCS data from a FITS file, it creates a FrameSet with this form; the base Frame represents pixel coordinates and the current Frame represents the FITS primary world coordinate system. The FrameSet may also contain other Frames representing any alternate axis descriptions stored in the FITS-WCS. This means that the FrameSet can be used as a Mapping from pixel coordinates to primary WCS, and can also be used as a Frame to determine the properties of the primary WCS.

The caller is free to select new base and/or current Frames at any time.

3.3.3. Integrity Restoration

Consider the simple case mentioned above where a FrameSet is used to describe the WCS in a 2-dimensional image. The FrameSet could for instance contain a pixel Frame as the base Frame and an (RA,Dec) Frame as the current Frame, connected together by a suitable Mapping. It is clearly possible to break the integrity of such a FrameSet, such that the Mapping no longer accurately describes the transformation from pixel coordinates to world coordinates. One obvious way in which this could be done is to change the current Frame so that it describes, say, galactic coordinates rather than (RA,Dec). The Mapping is left unchanged and so will still generate (RA,Dec) values, even though the FrameSet now claims that these are galactic coordinates[13].

In order to retain the integrity of the FrameSet, the Mapping must be replaced with one that generates the appropriate galactic coordinate values rather than (RA,Dec) values. One way in which this could be done is as follows, starting from the original unmodified FrameSet:

---

[12] The FrameSet class provides methods to set and retrieve the index of these two Frames.

[13] Note, Mappings are immutable and so the integrity of a FrameSet cannot be broken by making changes to a Mapping.

1. Create a copy of the current Frame (i.e., the (RA,Dec) Frame), and change the attributes of the copy so that it describes Galactic coordinates.

2. Use the `astConvert` method on the (RA,Dec) Frame to generate a Mapping from (RA,Dec) to galactic coordinates.

3. Add the galactic coordinates Frame into the FrameSet, using the above Mapping to connect it to the existing (RA,Dec) Frame.

4. Remove the original (RA,Dec) Frame.

The final FrameSet is unchanged in the sense that it still contains two Frames, but now the Mapping that connects them correctly generates the values described by the new current Frame (i.e., Galactic coordinates).

However, the above process is quite involved and prone to error, and so the FrameSet class itself provides generalised "integrity restoration" along the same lines, meaning that client code is relieved of the responsibility. Whenever any change is made to the current Frame within a FrameSet, the FrameSet class first takes a copy of the original Frame and makes the requested change to the copy. It then uses the `astConvert` method to generate a Mapping from the old Frame to the modified copy. It then adds the modified copy into the FrameSet, using the Mapping returned by *Convert* to connect it the original unmodified Frame. It finally removes the original original Frame.

In summary, the integrity restoration system within the FrameSet class means that whenever the properties of the current Frame are changed within a FrameSet, the Mappings within the FrameSet are automatically modified accordingly.

However there are times when this may not be what is wanted. For instance, if a mistake is made when setting the properties of a Frame prior to adding it into a FrameSet, then it should be possible to correct those mistakes without causing changes to be made to the Mappings within the FrameSet. This is possible by first getting a reference to the Frame, and then changing the Frame properties using that Frame reference, rather than the original FrameSet reference. Integrity restoration only occurs when the current Frame properties are changed via a FrameSet reference.

3.3.4. Searching a FrameSet for a Frame with Required Properties

As described earlier, the `astConvert` method defined by the Frame class attempts to find a Mapping between two arbitrary Frames[14]. Since the FrameSet class inherits the methods of the Frame class, the `astConvert` method can also be used on FrameSets. In this case the `astConvert` method will search through all the Frames in the FrameSet, starting with the current Frame, until a Frame is found for which a Mapping can be created. Since, in general, the Frames within a FrameSet all describe different domains, it is unlikely that more than one Frame

---

[14]This may or may not be possible depending on the nature of the two Frames.

will generate a Mapping, but the search order can be controlled by the caller in order to assign priority to specific Frames.

This allows code to search a FrameSet for a Frame that has specific properties. For instance, a Frame can be created with the required properties and then used as a "template" to search a FrameSet:

```
template = Ast.SkyFrame()
result = frameset.convert( template )
if result is None:
  print( "No celestial coord system found" )
```

In this example `frameset` is searched for a SkyFrame. If found, a new FrameSet is returned in which the base Frame is the matching Frame from `frameset`, the current Frame is a copy of `template` and the appropriate Mapping exists between them. In addition, the base Frame of `frameset` is set to indicate the matching Frame. If no matching Frame is found, a null reference is returned by *Convert*.

In the above example the specific celestial coordinate system represented by `template` was left unspecified when the SkyFrame was created. Consequently, the copy of `template` included in the `result` FrameSet returned by *Convert* inherits the celestial coordinate system of the matching Frame within `frameset`. If a specific system was specified when `template` was created, then that system would be given priority and be included in the returned `result` FrameSet.

3.4. Regions

Usually, axis values are stored simply as floating point values or arrays in the user's programming language and are not bound into AST objects. This is because the values are normally obtained and later used in this form (in other parts of the software, independently of AST). Also, speed of processing is vital. If axis values were routinely wrapped inside AST objects, the overhead of wrapping and unwrapping would be considerable, especially for small sets of positions that are processed repeatedly.

Nevertheless, axis values strictly only make sense within a particular coordinate system and this association can be made explicit, if required, by binding a set of positions to a specific Frame. This is performed by the Region class which encapsulates a Frame and a list of points, specified by their positions in the coordinate system of the Frame. The Region class inherits from the Frame class, so a Region can be used in place of the Frame that it contains.

The binding between the positions and the Frame is enforced by the Region class through "integrity restoration" similar to that described in section 3.3.3. If changes to the enclosed Frame are made through a Region reference and they result in a change to the Frame's coordinate system, then the coordinates stored in the Region are automatically transformed into that new coordinate system.

In practice, the Region class is abstract. It does not have a constructor function and is just a container class for other sub-classes which attach particular semantics to the set of positions that they contain. The simplest such sub-class is the PointList in which the positions are simple independent points. For example, a PointList might contain a set of star positions and a

SkyFrame indicating that these are in ICRS coordinates. If the SkyFrame's "System" attribute is changed via a Region reference to represent FK5 coordinates, then the stored coordinates would change so that the star positions remain fixed on the sky. As mentioned above, however, the PointList is not much used for general coordinate processing because of the overhead of packing and unpacking the object.

Other sub-classes of Region attach different semantics to the set of positions enclosed. For example, there are sub-classes to represent circles, ellipses, boxes, polygons. etc., where the enclosed positions may be fixed in number and define the Region's shape (the centre and one corner of a box, for example). These classes divide the Frame's domain into an "inside" and an "outside" and, as the "Region" name suggests, they are used to represent regions (intervals, areas, volumes, etc.) within domains. A negation method is provided that interchanges the inside and outside of a Region. Separate Regions may also be combined in pairs using boolean operations (AND, OR or XOR) and the CmpRegion class (which is analogous to the CmpMap and CmpFrame classes described elsewhere). This allows Regions of arbitrary complexity to be built out of simple components.

The significance of the "inside" and "outside" of a Region lies in its use as a Mapping, whose properties it inherits via the Frame class. Typically, a Region will map positions that lie inside it without change, but positions that lie outside are mapped to a special null value. This provides a test of whether a point is inside or outside a Region. A Prism class is also provided that allows Regions to be extruded into extra dimensions. For example, a 4-dimensional Prism can be formed by extruding a 2-dimensional Circle and a 2-dimensional Box. A 4-dimensional position is then regarded as inside the Prism if it lies inside the Circle on axes 1 and 2 and inside the Box on axes 3 and 4.

When combining Regions, it is not necessary for the coordinate systems represented by their enclosed Frames to be the same, but it must be possible to convert between them. If they differ, the astConvert method will be used to refer them both to a common coordinate system. Integrity restoration also means that any changes to a Region's Frame that changes the coordinate system will transform the boundary between the inside and outside of the Region into the new coordinate system. This may alter the nature of shapes such as circles, which might be distorted into ellipses or even more elaborate shapes under non-linear coordinate transformations. This is implemented by extending the enclosed Frame to become a FrameSet, in which the original undistorted coordinate system is preserved as the base Frame and the new coordinate system is the current Frame. When testing whether a point lies within the Region, it is first transformed from the current Frame into the base Frame of the Region's FrameSet and then tested against the Region as defined in that coordinate system.

In fact, Regions offer full support for enclosed FrameSets (as well as Frames) because a FrameSet is a sub-class of Frame and can therefore be supplied in place of a Frame when creating a Region. This allows a Region to be defined that has a known shape in one coordinate system but can be viewed through a range of other coordinate systems with corresponding changes

to its shape. This offers another means of creating Regions with novel shapes.

One obvious application of Regions is in the segmentation of datasets - for example, creating outlines that enclose objects or other regions of interest in an image. However, by using Regions as Mappings and combining them with other Mappings, it is possible to limit the range of coordinates that the Mapping will accept, such that null values are returned outside the Region. This effectively "clips" the range of validity of the Mapping. This clipping can be arranged to occur in any coordinate system: either the input or output space of the Mapping, an intermediate space, or one devised purely for the purpose of clipping. These techniques find particular use in graphical applications for limiting the extent of drawing operations.

4. Serialisation, and FITS-WCS

AST includes a set of *Channel* classes, which allow AST objects to be serialised in various ways for persistent external storage. The basic Channel class has a method that converts an in-memory AST object into a set of text strings. By default these text strings are simply written to standard output, but an external "sink" function can be supplied to the Channel constructor which — if supplied — will be called by the Channel class to store each text string in some associated external store. The Channel class has another method that does the inverse — it calls a supplied "source" function to read a set of text strings from an external store, and then creates an in-memory AST object from the text strings. A round-trip (write followed read) is lossless. The textual descriptions of AST objects produced by the Channel class use a bespoke block structured format specific to AST.

As an example, the following Python code:

```
map1 = starlink.Ast.PermMap( [3,1],
                             [2,-1,1], 12.2 )
map2 = starlink.Ast.ZoomMap( 3, 4.0 )
cmpmap = starlink.Ast.CmpMap( map1, map2 )

channel = Ast.Channel()
channel.write( cmpmap )
```

produces the following output[15]:

```
Begin CmpMap    # Compound Mapping
   Nin = 2      # Number of input coordinates
   Nout = 3     # Number of output coordinates
IsA Mapping     # Mapping between coordinate systems
   MapA =       # First component Mapping
      Begin PermMap   # Coordinate permutation
         Nin = 2      # Number of input coordinates
         Nout = 3     # Number of output coordinates
      IsA Mapping     # Mapping between coordinate systems
         Out1 = 2     # Output coordinate 1 = input coordinate 2
         Out2 = -1    # Output coordinate 2 = constant no. 1
         Out3 = 1     # Output coordinate 3 = input coordinate 1
         In1 = 3      # Input coordinate 1 = output coordinate 3
         In2 = 1      # Input coordinate 2 = output coordinate 1
         Nconst = 1   # Number of constants
         Con1 = 12.2  # Constant number 1
      End PermMap
   MapB =       # Second component Mapping
```

---

[15]Within the Python interface, the *print* function uses a Channel to produce a listing of the given AST object. So the same output could more simply be produced by "print(cmpmap)".

12

```
        Begin ZoomMap    # Zoom about the origin
            Nin = 3      # Number of input coordinates
        IsA Mapping      # Mapping between coordinate systems
            Zoom = 4     # Zoom factor
        End ZoomMap
End CmpMap
```

Attributes of the Channel class can be used to enable or disable inclusion of comments, indentation, defaulted values, *etc*.

Sub-classes of Channel are provided that can encode this lossless block-structured object description into an XML format (the XmlChan class), or into a set of FITS header cards (the FitsChan class).

4.1. Foreign Encodings

In addition to the "native" block-structured object description outlined above, AST can also read and write a range of "foreign" encodings that allow data exchange with other systems. Because this is essentially a format-conversion process it is not usually lossless. AST's flexibility allows it to represent information that it may not be possible to store in an external format, so some data loss may occur when writing and only particular types of AST object can be successfully converted. Conversely, when reading, there is normally little need to lose information, although the more obscure features of some formats may not all be supported.

4.1.1. FITS Encodings

There are a number of conventions (unrelated to AST) for storing WCS information in FITS header cards and the rules governing their use is complex. Most programmers are unlikely to understand these technical details, so AST wraps them up in a simple high-level model provided by the FitsChan class. This provides a read/write interface between AST objects and FITS header cards (where each card represents a "keyword=value" pair plus an optional comment).

The FitsChan class currently provides support for almost all of papers I, II and III in the FITS-WCS series (Calabretta and Greisen, 2002; Greisen et al., 2006; Greisen and Calabretta, 2002) [16]. As yet it does not provide support for converting between the AST TimeFrame class and the FITS-WCS description of time axes (Rots et al., 2015). It also supports a variety of unofficial encodings that have been devised by particular projects over many years, including several popular methods for describing focal plane distortion.

Normally, FITS headers do not simply hold WCS information; they also contain many additional cards describing other aspects of a dataset. Care is therefore needed to ensure these other cards are not accidentally removed or over-written while processing WCS information. Also, the various FITS encodings of WCS information may overlap in their use of FITS keywords, presenting an opportunity for keyword clashes.

To handle this, a FitsChan contains a buffer which may be loaded with FITS header cards, normally read directly from a data file. The FitsChan analyses these cards, determines how any contained WCS information has been stored and makes this available through its "Encoding" attribute. This takes values such as "FITS-WCS", "FITS-IRAF", "NATIVE ", etc. to indicate which convention has been used to store the information. A subsequent call to the FitsChan's astRead method will then perform a "destructive read" on the buffer of FITS cards. This reads the AST object that they describe and then deletes all the relevant FITS cards. The method returns a FrameSet which holds the WCS information (see section 3.3.2) and at the same time this information is removed from the FITS header buffer, leaving it clear for new WCS information to be written without the possibility of any keyword clashes.

To write a FrameSet[17] to a FITS header, a FitsChan's astWrite method is used. This uses the "Encoding" attribute to determine which FITS header convention to use and then adds the resulting FITS header cards to the FitsChan buffer, from where they may be transferred to an output file. If changes have been made to a FrameSet such that it cannot be stored using the chosen FITS encoding, then the astWrite method will fail. AST will go to some lengths to simplify the information so that it is compatible with the data model of the selected encoding, but that is not always possible. In such cases, an alternative encoding can be chosen simply by modifying the "Encoding" attribute of the FitsChan.

4.1.2 The STC-S Encoding The StcsChan class provides an interface between AST and space-time metadata from a subset of the "STC-S" scheme (Rots, 2011) developed by the International Virtual Observatory Alliance (IVOA). This allows AST Region objects to be converted to and from textual descriptions that use the STC-S conventions .

5. Fields of Application

This section describes some of the ways in which AST can be used.

Whilst the AST library provides many facilities that are useful when describing and using the WCS information attached to a data array, it is not limited to that field. Its generalised design enables it to be used in any situation where relationships between several different coordinate systems need to be managed. Just to emphasise that point, WCS handling is included as the last item in this section.

5.1. Generalised Plotting

Most graphics systems provide many coordinate systems through which the display can be addressed - device pixel coordinates, normalised viewport coordinates, *etc*. However, since the details of the relationships between such coordinate systems may vary from system to system, AST does not attempt to describe them all. Instead it just assumes there is at least one such coordinate system available and provides facilities that allow it to be attached to that system and to use it to perform plotting[18].

---

[16]Support for the tabular format associated with the "-TAB" code described in FITS-WCS paper III is limited to 1-dimensional(*i.e.* separable) axes.

[17]A FrameSet is the only valid object that can be written to a FitsChan.

[18]The user is still free to use the underlying graphics system directly when required.

In addition, AST allows application software to define extra — potentially non-linear — coordinate systems, and to integrate them with the coordinate system provided by the underlying graphics package. For this purpose AST provides two classes — *Plot*, which provides two-dimensional plotting facilities, and *Plot3D*, which provides three-dimensional plotting facilities.

5.1.1. Use of External Plotting Packages

The plotting classes within AST do not themselves include any facilities for placing "ink onto paper" — an external graphic library must be supplied to draw graphical primitives such as straight lines, markers and character strings. The plotting classes within AST will then use the primitive facilities of this underlying graphics system to draw more complex entities such as annotated coordinate grids, *etc*.

This separation between coordinate handling and drawing enables the sophisticated plotting capabilities of AST to be used within many different systems and languages[19].

Each of the underlying graphics systems (the 2-d and 3-d plotting classes use separate systems) can be specified in two ways:

1. At build-time. In this method, a module must be supplied (callable from C) that provides implementations of a set of wrapper functions that AST uses to perform primitive drawing operations. Each such wrapper function makes appropriate calls to the underlying graphics system to perform its work. This module is linked into the executable at build-time, in place of the default module provided by AST.

2. At run-time. In this method, application code registers pointers to the graphics wrapper functions with AST, whilst the executable is running. The registered function pointers are used in preference to any functions specified at build-time.

The AST library includes modules that allow the PGPLOT graphics package (Pearson, 1991, ascl:1103.002) to be used for drawing by both the Plot class and the Plot3D class (see the following sections).

5.1.2. Two Dimensional Plotting

A *Plot* is a subclass of FrameSet and can therefore be considered to be a FrameSet 'with some extra facilities". A Plot does not represent the graphical content itself, but is a route through which plotting operations, such as drawing lines and curves, are conveyed on to a plotting surface to appear as visible graphics.

When considered as a FrameSet, the base Frame within a Plot corresponds to the Cartesian coordinate system used to specify positions to the underlying graphics system[20]. The bounds of the plotting area within this coordinate system must be specified when the Plot is created. A typical Plot may for instance have

a base Frame that describes millimetres from the bottom left corner of the plotting area.

The current Frame within a Plot corresponds to "user" coordinates — *i.e.*, the coordinate system in which the application code wishes to specify positions. Since the Plot is a form of FrameSet, it will also include the Mappings needed to transform "user" coordinates into the corresponding graphics coordinates. A typical Plot may for instance have a current Frame that describes (RA,Dec) positions on the sky.

When a Plot is created, an existing FrameSet is supplied together with the bounds of a box within its base Frame (typically image pixel coordinates). A Mapping (either linear or logarithmic) is created that maps this box onto a specified area within the graphics coordinates systems. The Plot constructor then initialises the new Plot to be a copy of the supplied FrameSet, and adds in a new Frame describing the graphics coordinates system, using the above Mapping to connect the new Frame to the base Frame. This new Frame is then made the base Frame in the Plot. In effect, the Plot becomes attached to the plotting surface in rather the same way that a basic FrameSet might be attached to (say) an image.

The Plot class has methods that can draw markers, geodesic curves, text strings, *etc*. The application code supplies the positions of these objects within the "user" coordinate system (*i.e.*, the current Frame of the Plot — *e.g.*, (RA,Dec)). The Plot class then uses the Mappings stored within the Plot to transform them into the graphics coordinate system (*i.e.*, the base Frame of the Plot), before invoking the appropriate wrapper functions to instruct the underlying graphics system to draw the required primitives. Whilst this is a fairly simple process for items such as graphical markers that only have one associated coordinate (*i.e.*, the centre of the marker), it is a more complex process for items such as geodesic curves that span a wide range of different coordinates. In this case, the curve is represented by a set of positions spread along the curve in user coordinates. All these positions are transformed into graphics coordinates before being plotted in the form of a "poly-line". The density of points varies along this line, and is chosen to ensure that any discontinuities or highly non-linear sections of the curve are drawn with sufficient accuracy. All aspects of the generated plot can be controlled via attributes of the Plot class.

The most comprehensive drawing method supplied by the Plot class produces a complete set of annotated axes describing the area of user coordinates visible within the plotting area. Some examples are shown in Fig. 3.

5.1.3. Three Dimensional Plotting

The *Plot3D* class extends the Plot class to provide plotting facilities in three dimensions. The basic model is the same as for the Plot class — a Plot3D object is a FrameSet in which the base Frame represents the 3-dimensional Cartesian coordinate system used by the underlying 3-dimensional plotting package, and the current Frame represents 3-dimensional "user" coordinates.

The projection from 3-dimensional graphics coordinates onto a 2-dimensional plotting surface is handled by the underlying graphics package. This may be surprising, given that AST

---

[19]It is known that AST graphics have been used with plotting packages written in Java, Perl, Tk-Tcl and Python, as well as C and Fortran.

[20]This coordinate system must have been defined previously using appropriate calls to the underlying graphics system.
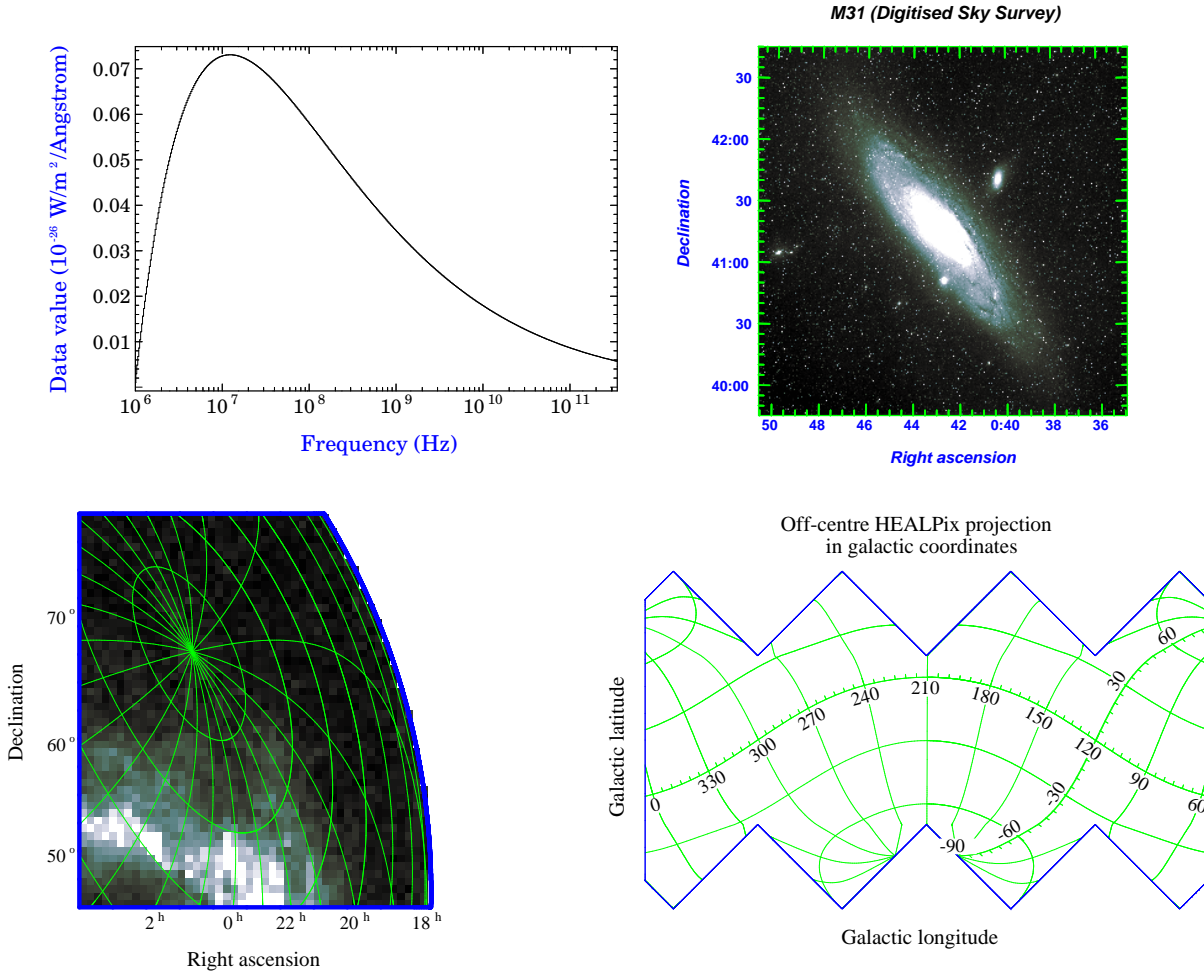
Figure 3: A selection of annotated axes created by the Plot class. Note, Plot only creates the annotated axes — the background images were produced using PGPLOT directly. The PGPLOT library was used as the underlying graphics package for these plots.

could itself handle this sort of projection. But the decision was taken so that Plot3D graphics wrapper functions could use the full range of features offered by a dedicated 3-dimensional plotting packages, rather than restrict it to the more limited features offered by a typical 2-dimensional package.

Fig. 4 shows an example of the 3-dimensional annotated axes produced by Plot3D. The Plot3D class was originally developed to support 3-D coordinate grid plotting in the Starlink GAIA display tool (Draper et al., 2008, ascl:1403.024) using the VTK 3-d graphics package (Hanwell et al., 2015).

5.2. Flux and Data Unit Transformation

In the majority of scientific data sets, each data point has a *position* and a *value*. Whilst the concepts described in this paper may seem more naturally associated with the positional information, they can also be applied to the value (or values) associated with each data point. A position within some N-dimensional space is specified by a set of N axis values. Likewise, a value can also be represented by a set of axis values. In

the majority of common cases data values are 1-dimensional - for instance a temperature, or a sky brightness. But there are also some common multi-dimensional cases such as Stokes vector (*I*, *Q*, *U*, *V*).

The functional division into Mappings, Frames and Frame-Sets used throughout AST is equally applicable to the problem of describing data *values* as it is to describing data *positions*. To do so requires a set of specialised Frame classes to be created to describe each set of equivalent data value systems. For instance a hypothetical *TemperatureFrame* could be created that encapsulates the meta-data needed to transform between different temperature scales (Celsius, Fahrenheit, Kelvin, *etc*).

Currently, provision of such specialised Frames within AST is limited, but this may change in future:

FluxFrame : Supports 1-dimensional axes that represent the following astronomical flux systems:

  • Flux per unit frequency ($W/m^2/Hz$)

15

Figure 4: Annotated axes created by the Plot3D class. The 3D interface to the PGPLOT library included with AST was used as the underlying graphics package for these plots.

- Flux per unit wavelength ($W/m^2/Angstrom$)

- Surface brightness in frequency units ($W/m^2/Hz/arcmin^2$)

- Surface brightness in wavelength units ($W/m^2/Angstrom/arcmin^2$)

Any dimensionally equivalent units can be used in place of the default units listed above. All flux values are assumed to be measured at the same frequency or wavelength (specified by an attribute of the FluxFrame). Thus this class is more appropriate for use with images than spectra.

Frame : The basic Frame class can be used to represent single-system data value axes (*i.e.* Frames that know only one system for specifying positions within its domain), using any suitable system of units. Such Frames can be used to convert axis values into any dimensionally equivalent set of units. For instance, a basic Frame used to represent stellar mass could be set to use Kg, Solar mass, *etc.* as its units, and will automatically modify the associated Mappings each time the units are changed (assuming the Frame is part of a FrameSet).

5.3. Attaching WCS Information to Datasets

The most common reason for using AST is to allow the position of one or more data values to be described in a range of possible coordinate systems. For this purpose, a serialised FrameSet is often stored with the data on the assumption that the "natural coordinate system" of the data structure corresponds to a known Frame (usually the base Frame) in the FrameSet. The most common data structure is an N-dimensional regular grid

of values, in which case the base Frame in the FrameSet is assumed to describe "pixel coordinates" within the grid (generalising the term "pixel" to data of any dimensionality). Different systems exist for enumerating the pixels within an array - for instance, some start counting at zero and some at one. AST allows multiple pixel coordinate systems to be described and used within a single FrameSet, assuming that the Mappings between such Frames have been set up and incorporated into the FrameSet correctly.

The following sections describe just a few of the more common WCS operations that may be achieved using AST:

5.3.1. Validating WCS

The generalised description of WCS provided by AST makes it possible to write applications in a domain-agnostic manner. For instance, an application that uses WCS to align two data sets need not know whether the data sets are two-dimensional images of the sky, one dimensional spectra, 3-dimensional time-space cubes, *etc.* However, this will not always be possible, and it will still often be the case that an applications needs to check that the WCS in the supplied data is consistent with the specific requirements of the application. The `astFindFrame` method can often be used for this purpose. The application creates a Frame that acts as a template for the coordinate systems required by the application, and passes this Frame, together with the WCS FrameSet read from the data, to the `astFindFrame` method, which then searches the FrameSet looking for a Frame that can be matched against the template. If such a Frame is found, information about the specific Frame found is returned, together with a Mapping that connects that Frame to one of the Frames in the WCS FrameSet.

As a trivial example, consider an application that requires 1-dimensional data but does not care what that one dimension represents, it creates a 1-dimensional basic Frame to act as a template, leaving all Frame attributes at their default values in the template (default values act as wild-cards during the searching process). The application then passes this template Frame, together with the WCS FrameSet read from the data file, to `astFindFrame`. Each Frame in the WCS FrameSet is then compared to the template to see if a match is possible. In this case, a basic Frame with no set attributes will match any 1-dimensional Frame of any form, but will not match Frames with more than one dimension.

A specialised application for processing spectral data cubes may use for its template a 3-dimensional CmpFrame (compound Frame) containing a 2-dimensional SkyFrame (celestial longitude/latitude axes) and 1-dimensional SpecFrame (a spectral axis). The `astFindFrame` will then only match 3-dimensional Frames with similar properties. Note, if the System attribute is set to specific values in the SkyFrame and/or SpecFrame contained within the template, then the Mapping returned by `astFindFrame` will include the transformations needed to convert from the system of the matching Frame to that of the template Frame. So for instance, if the template axes are ($RA, Dec, Wavelength$) and the WCS FrameSet contains a CmpFrame with ($Frequency, GalacticLongitude, GalacticLatitude$) axes,

16

the Mapping returned by `astFindFrame` will include an axis permutation, and the transformations needed to convert from frequency to wavelength, and from Galactic coordinates to (*RA*, *Dec*).

In other words, `astFindFrame` provides arbitration between the WCS requirements of the application, and the WCS information that is available in the supplied data set, returning a null result if the requirements of the application are not met by the data set.

5.3.2. Merging WCS information (Alignment)

If a data processing operation involves combining two or more datasets in some way, it will usually be necessary to form a connection between the WCS in the two datasets, so that corresponding elements in each dataset can be identified. This process can be considered one of alignment. For example, in the case of images of the sky, we may wish to align them prior to combining them into a mosaic. If each dataset is calibrated using a FrameSet whose base Frame corresponds with the underlying data elements (e.g. pixels), then alignment involves finding a transformation between the base Frames of the two FrameSets involved.

This operation is performed using the astConvert method introduced in section 3.2.2, whose role is to find a transformation between the coordinate systems represented by two Frames. Because FrameSets are also Frames, astConvert may be used to convert between two FrameSets. In this case, however, there is a choice of how to find a transformation between the two Frames that they represent (their current Frames). This is because links between two FrameSets can potentially be formed in a number of places by matching Frames within each FrameSet that have the same domain.

An example should make this clear. Suppose we have two images, each calibrated with FrameSets whose base Frames represent image pixel coordinates. Also suppose that each FrameSet contains two additional Frames that give the image's position in the focal plane of a telescope (mm) and its position on the sky (RA,Dec). If we want to align these images, we can do it in three ways by matching: (1) corresponding pixels, (2) corresponding focal plane positions or (3) corresponding positions on the sky. Any of these is a legitimate way to align the two images and may be appropriate for different purposes.

When presented with the two FrameSets,[21] astConvert has the same choices. This ambiguity may be removed by supplying the name of the domain that we wish to use for alignment. In this example, we might select "SKY" and astConvert would then try to find a route between the two FrameSets that joins a pair of SkyFrames (representing celestial coordinates), one from each FrameSet. A domain "search path" may also be given, in which case each listed domain will be tried in turn until one succeeds.

The result returned by astConvert will be a FrameSet which, when used as a Mapping, converts pixel coordinates in one image into pixel coordinates in the second image, such that the two images are aligned in the selected domain.

5.3.3. Modifying WCS Information

If an application creates an output data array by applying some geometrical transformation to an input data array, it should also store appropriately modified WCS in the output. The concatenation and simplification of Mappings described in section 3.1 makes this easy. The application should create a Mapping that describes the geometric transformation that has been applied to the pixel array, and then "re-map" the pixel Frame within the WCS FrameSet read from the input data set. The resulting modified FrameSet should be stored in the output data set. Re-mapping a Frame within a FrameSet means appending the supplied Mapping to the existing Mapping that connects the Frame with its parent in the FrameSet. A method is provided to do this, which also simplifies the resulting compound Mapping if possible.

For instance, if an application rotates an input image to create an output image, the application should create a 2-dimensional MatrixMap to describe the rotation[22] and then invoke astRemapFrame, supplying the MatrixMap and the WCS FrameSet from the input image. The modified FrameSet would then be stored in the output image.

5.3.4. Using WCS for Data Resampling and Regridding

The previous section described how to modify the WCS to take account of a geometric transformation of a data set, assuming a Mapping describing the transformation is available. If such a Mapping is available, it can also be used to perform the actual resampling or regridding of the pixel values themselves. AST provides several methods that will perform such resampling or regridding, using any of a wide range of alternative sampling kernels. Alternatively, an externally defined sampling kernel can be used.

The benefit of using these methods within AST rather than simply transforming every pixel position within the application code, is that the AST methods attempt to speed up the operation by using linear approximations to the supplied transformation if possible. Data sets are constantly increasing in size, and transforming every pixel position within a large data set using a long and complicated transformation can be an expensive operation. The AST methods divide the input data set in half along each axis, and find a linear approximation to the Mapping over each resulting quadrant. If one of these approximations meets a user-specified accuracy, the approximation is used to transform all pixel positions within its quadrant. Any quadrants that do not meet the accuracy requirements are subdivided again and new approximations are found for each of the new sub-quadrants. This process repeats recursively until the sub-quadrants become so small that there is likely to be little time-saving in sub-dividing them any further. At this point the full transformation is used on all pixels within any such sub-quadrants.

---

[21]Because astConvert acts on the current Frame of a FrameSet, the two FrameSets must actually be temporarily modified to interchange their base and current Frames. Inverting them will have this effect.

[22]A pair of ShiftMaps will also be needed if the rotation is not around the pixel origin.

6. Things we Would do Differently Now

6.1. Language Choice

There was no C++ standard in 1996 when AST development commenced. So for portability reasons we chose to write AST in ANSI C, developing our own infrastructure to support object-orientation. Things are different now — C++ is much more mature with greater standardisation. For that reason we would choose C++ over C if starting AST development now[23].

6.2. Architecture

The current architecture of AST makes it difficult for non-Starlink developers to create optional extensions or plug-ins to AST that provide added value. This was seen initially as a positive feature that would encourage standardisation, and avoid proliferation of multiple incompatible versions of AST object format. However, with hind sight the monolithic nature of the library has turned out to be a barrier to wider adoption. Potential users often want to use only some small part of AST's wide functionality, and feel daunted by a large library and a fat document covering many things that are irrelevant to them. Also, the custom infra-structure used for object orientation within AST is often seen as hard to learn. Some simple and well documented means for an external library to use this infrastructure would have been useful.

6.3. Floating-point precision

The restriction that all coordinate values are represented by double precision values may be a problem for time axes, which can sometimes require quad precision.

6.4. Coordinate systems versus Domains

The distinction between a coordinate system and a domain has caused confusion and misunderstandings on several occasions. Coordinate systems are mathematical abstractions of various types (Cartesian, polar, *etc*). Coordinate system are used to describe positions within a physical space (*i.e.* a domain). In this sense a domain can encapsulate several alternative coordinate systems, any of which can be used to describe positions in the domain. With hind-sight, it may have been useful to make this distinction clearer by separating coordinate systems and domains into different classes.

6.5. Angle units

Our decision to use radians rather than degrees to represent angle values has proved unpopular.

7. The AST library

The AST library has been developed over a number of years (Berry and Draper, 2010; Berry, 2001, 2004, 2008; Berry and Jenness, 2012; Warren-Smith and Berry, 1998, 2000) and is written in C with no dependencies. It includes code from WCSLIB (Calabretta, 2006, ascl:1108.003), PAL (Jenness and Berry, 2013) and SOFA (Hohenkerk, 2011, ascl:1403.026) but does not depend on those libraries. There are language bindings for Fortran, Java, Perl and Python.

8. Acknowledgements

References

Astropy Collaboration, 2013. Astropy: A community Python package for astronomy. Astron Astrophys 558, A33. doi:10.1051/0004-6361/201322068, arXiv:1307.6212.

Berry, D., Draper, P., 2010. Using the AST Library to Create and Use STC-S Region Descriptions, in: Mizumoto, Y., Morita, K.I., Ohishi, M. (Eds.), Astronomical Data Analysis Software and Systems XIX, volume 434 of *ASP Conf. Ser.*. p. 213.

Berry, D.S., 2001. Providing Improved WCS Facilities Through the Starlink AST and NDF Libraries, in: Harnden, Jr., F.R., Primini, F.A., Payne, H.E. (Eds.), Astronomical Data Analysis Software and Systems X, volume 238 of *ASP Conf. Ser.*. p. 129.

Berry, D.S., 2004. Developments in the Starlink AST Library - an Intelligent WCS Management System, in: Ochsenbein, F., Allen, M.G., Egret, D. (Eds.), Astronomical Data Analysis Software and Systems (ADASS) XIII, volume 314 of *ASP Conf. Ser.*. p. 412.

Berry, D.S., 2008. Developments in the AST Library, in: Argyle, R.W., Bunclark, P.S., Lewis, J.R. (Eds.), Astronomical Data Analysis Software and Systems XVII, volume 394 of *ASP Conf. Ser.*. p. 635.

Berry, D.S., Jenness, T., 2012. New Features in AST: A WCS Management and Manipulation Library, in: Ballester, P., Egret, D., Lorente, N.P.F. (Eds.), Astronomical Data Analysis Software and Systems XXI, volume 461 of *ASP Conf. Ser.*. p. 825.

Calabretta, M.R., 2006. Implementation of -TAB coordinates in WCSLIB, in: Gabriel, C., Arviset, C., Ponz, D., Enrique, S. (Eds.), Astronomical Data Analysis Software and Systems XV, volume 351 of *ASP Conf. Ser.*. p. 591.

Calabretta, M.R., Greisen, E.W., 2002. Representations of celestial coordinates in FITS. Astron Astrophys 395. doi:10.1051/0004-6361:20021327, arXiv:astro-ph/0207413.

Currie, M.J., Berry, D.S., Jenness, T., Gibb, A.G., Bell, G.S., Draper, P.W., 2014. Starlink Software in 2013, in: Manset, N., Forshay, P. (Eds.), Astronomical Data Analysis Software and Systems XXIII, volume 485 of *ASP Conf. Ser.*. p. 391.

Draper, P., 1998. New Features in GAIA. Starlink Bulletin 20, 7.

Draper, P.W., Berry, D.S., Jenness, T., Economou, F., Currie, M.J., 2008. GAIA-3D: Volume Visualization of Data-Cubes, in: Argyle, R.W., Bunclark, P.S., Lewis, J.R. (Eds.), Astronomical Data Analysis Software and Systems XVII, volume 394 of *ASP Conf. Ser.*. p. 339.

Eaton, N., McIlwrath, B., 2013. AGI — Applications Graphics Interface Library. Starlink User Note 48, Starlink Project.

---

[23]Writing in higher level languages such as Java or Python would make it more difficult to write interfaces for other languages.

[24]https://github.com/Starlink/starlink

Evans, I. (Ed.), 2013. Twenty Years of ADASS. volume MP-006 of *ASP Monograph*. Astronomical Society of the Pacific. ISBN: 978-1-58381-822-0.

Greisen, E.G., Calabretta, M.R., Valdes, F.G., Allen, S.L., 2006. Representations of spectral coordinates in FITS. Astron Astrophys 446. doi:10.1051/0004-6361:20053818, arXiv:astro-ph/0507293.

Greisen, E.W., Calabretta, M., 1995. Representations of Celestial Coordinates in FITS, in: Shaw, R.A., Payne, H.E., Hayes, J.J.E. (Eds.), Astronomical Data Analysis Software and Systems IV, volume 77 of *ASP Conf. Ser.*. p. 233.

Greisen, E.W., Calabretta, M.R., 2002. Representations of world coordinates in FITS. Astron Astrophys 395. doi:10.1051/0004-6361:20021326, arXiv:astro-ph/0207407.

Hanwell, M.D., Martin, K.M., Chaudhary, A., Avila, L.S., 2015. The Visualization Toolkit (VTK): Rewriting the rendering code for modern graphics cards. SoftwareX 1–2, 9–12. doi:10.1016/j.softx.2015.04.001.

Hohenkerk, C., 2011. Standards of Fundamental Astronomy. Scholarpedia 6, 11404. doi:10.4249/scholarpedia.11404.

Holub, A.I., 1992. Programming with Objects in C and C++. McGraw-Hill. ISBN 0-07-029662-6.

Jenness, T., 2015. Reimplementing the Hierarchical Data System using HDF5. Astron. Comp. 12, 221–228. doi:10.1016/j.ascom.2015.02.003, arXiv:1502.04029.

Jenness, T., et al., 2015. Learning from 25 years of the extensible *N*-Dimensional Data Format. Astron. Comp. 12, 146–161. doi:10.1016/j.ascom.2014.11.001, arXiv:1410.7513.

Jenness, T., Berry, D.S., 2013. PAL: A Positional Astronomy Library, in: Friedel, D.N. (Ed.), Astronomical Data Analysis Software and Systems XXII, volume 475 of *ASP Conf. Ser.*. p. 307.

Joye, W.A., Mandel, E., 2003. New Features of SAOImage DS9, in: Payne, H.E., Jedrzejewski, R.I., Hook, R.N. (Eds.), Astronomical Data Analysis Software and Systems XII, volume 295 of *ASP Conf. Ser.*. p. 489.

Lawden, M., 1989. Software and documentation news. Starlink Bulletin 4, 7.

Lawden, M., 1998. New Products. Starlink Bulletin 20, 6.

Lupton, W.F., 1989. HDS — Hierarchical Data System. Starlink System Note 27, Starlink Project.

Mink, J., Mann, R.G., Hanisch, R., Rots, A., Seaman, R., Jenness, T., Thomas, B., 2015. The Past, Present and Future of Astronomical Data Formats, in: Taylor, A.R., Rosolowsky, E. (Eds.), Astronomical Data Analysis Software and Systems XXIV, volume 495 of *ASP Conf. Ser.*. pp. 11–18. arXiv:1411.0996.

Pearson, T.J., 1991. Caltech VLBI Analysis Programs, California Institute of Technology. Bulletin of the American Astronomical Society 23, 991–992.

Rots, A.H., 2011. Space-time coordinate metadata for the virtual observatory version 1.33. IVOA Recommendation. arXiv:1110.0504.

Rots, A.H., Bunclark, P.S., Calabretta, M.R., Allen, S.L., Manchester, R.N., Thompson, W.T., 2015. Representations of time coordinates in FITS. Time and relative dimension in space. Astron Astrophys 574, A36. doi:10.1051/0004-6361/201424653, arXiv:1409.7583.

Thomas, B., et al., 2015. Learning from FITS: Limitations in use in modern astronomical research. Astron. Comp. 12, 133–145. doi:10.1016/j.ascom.2015.01.009, arXiv:1502.00996.

Wallace, P.T., 1994. The SLALIB Library, in: Crabtree, D.R., Hanisch, R.J., Barnes, J. (Eds.), Astronomical Data Analysis Software and Systems III, volume 61 of *ASP Conf. Ser.*. p. 481.

Warren-Smith, R.F., 1989. TRANSFORM —Coordinate Transformation Facility. Starlink User Note 61, Starlink Project.

Warren-Smith, R.F., Berry, D.S., 1998. World Coordinate Systems as Objects, in: Albrecht, R., Hook, R.N., Bushouse, H.A. (Eds.), Astronomical Data Analysis Software and Systems VII, volume 145 of *ASP Conf. Ser.*. p. 41.

Warren-Smith, R.F., Berry, D.S., 2000. Recent Developments to the AST Astrometry Library, in: Manset, N., Veillet, C., Crabtree, D. (Eds.), Astronomical Data Analysis Software and Systems IX, volume 216 of *ASP Conf. Ser.*. p. 506.

Warren-Smith, R.F., Berry, D.S., 2013. AST – A Library for Handling World Coordinate Systems in Astronomy. Starlink User Note 211, Starlink Project, STFC.

Warren-Smith, R.F., Lawden, M.D., McIlwrath, B.K., Jenness, T., Draper, P.W., 2008. HDS — Hierarchical Data System. Starlink User Note 92, Starlink Project.

Wells, D.C., Greisen, E.W., Harten, R.H., 1981. FITS: a flexible image transport system. Astron Astrophys Supp 44, 363–370.